Wei Lin

(213)458-8863 | wlin374@gatech.edu | www.linkedin.com/in/weilin-aie | No Sponsorship Required

EDUCATION

GEORGIA INSTITUTE OF TECHNOLOGY Atlanta, Georgia
MS in Computer Science, AI Track Expected May 2027

PROFESSIONAL EXPERIENCE

GenShinChon International Trading Co., Ltd Kowloon Bay, Hong Kong
ML Engineer Intern Jan 2024 - Mar 2024

- Automated vendor specification parsing using **Java** and **Docker** to standardize data ingestion pipelines, eliminating manual entry and reducing quote generation latency by **18%**.
- Deployed **PostgreSQL** on **AWS** within **Kubernetes** clusters to manage high-concurrency lookups for model training data, achieving **15% faster response times** via index optimization.
- Integrated **OpenAI API** into automated CI/CD workflows to build semantic information retrieval engine for machinery manuals, enhancing document search accuracy by **20%**.

Ipro Financial Press Ltd Central, Hong Kong
Backend Engineering Intern Mar 2023 - May 2023

- Built containerized microservices using **Docker** to deploy LLM proofreading models on **AWS EKS**, reducing manual financial document review turnaround time by 15%.
- Implemented distributed validation layer using **TypeScript** and **Amazon SQS** to sanitize model-extracted outputs, decreasing schema mapping anomalies and parsing failures by 18%.
- Optimized **MongoDB** query performance and **RESTful API** execution within **Kubernetes** clusters, utilizing Horizontal Pod Autoscaling (HPA) to deliver **12% latency reduction** during peak IPO submission windows.

PROJECT AND RESEARCH EXPERIENCE

Aura Travel LLC Troy, Michigan
AI Engineer & Founder
Python | LangChain | ChromaDB Feb 2026 - Present

- Developed RAG pipelines using LangChain, semantic chunking, and vector-based ANN retrieval, improving context relevance by 15%.
- Implemented travel report generation workflows utilizing **Chain-of-Thought (CoT)** and **prompt chaining**, resulting in **12% reduction** in model hallucinations.
- **Built** orchestration-worker multi-agent systems via LangGraph and tool calling, resulting in **18% decrease** in end-to-end latency for itinerary generation.

Google AI: "AI Socrates" Learning Agent
SWIFT | MVVM |Gemini API Dec 2025 - Feb 2026

- Built **MVVM** architecture using **Swift** to manage asynchronous **Gemini API** data streaming. Employed state machine workflows resolving race conditions across 4 learning phases, preventing UI thread freezing and improving interaction latency by **15%**.
- Automated native iOS build deployments and unit testing workflows utilizing **Harness CI** platforms. Streamlined continuous integration pipelines resolving manual compilation bottlenecks, accelerating prototype feature iterations and reducing version release turnaround times by **20%**.

Mini Open WebUI | FullStack AI Chat Application
FastAPI | SvelteKit | SQLite | RAG May 2026 - Jun 2026

- Built FastAPI backend with SQLite persistence, session authentication, and per-user access controls for conversations, messages, and documents.
- Implemented **SvelteKit** interface across chat, document, and settings views, integrating streaming responses, document uploads, model configuration, and loading/error states.
- Developed RAG pipeline combining document chunking, keyword retrieval, context augmentation, and retrieved-source display through OpenAI-compatible APIs.

Technical Skills

•**Languages & Frameworks**: Python, Java, C++, JavaScript, TypeScript, SQL, Spring Boot, JUnit, RESTful APIs
•**Cloud & DevOps**: AWS (Lambda, S3, EC2, Athena, Redshift, ECS), Docker, Kubernetes, Jenkins, Terraform, CI/CD
•**ML/AI**: LangGraph, LangChain, RAG, PGVector, OpenAI APIs, QLoRA, Vision Transformers, Active Learning